

# Operating System Labs

Yuanbin Wu  
cs@ecnu

# Announcement

- Project 0 due
  - 21:00, Sep. 19

# Operating System Labs

- Introduction of I/O operations
- Project 0b
  - Sorting

# Operating System Labs

- Manipulate I/O
  - System call
    - File descriptor
    - No buffering
  - Standard library
    - FILE object
    - Buffering

# Operating System Labs

- Manipulate I/O
  - System call
    - File descriptor
  - Standard library
    - FILE object
    - Buffer/non-buffer

# I/O System Calls

- 5 basic system calls
  - `open()`, `read()`, `write()`, `lseek()`, `close()`
- I/O without buffering
- File sharing
  - understand file descriptor
  - `dup()` `dup2()`
- Other
  - `fcntl()`, `sync()`, `fsync()`, `ioctl()`

# File Descriptor

- File descriptor
  - Allocated when open a file
  - “ID” of the file **in the process** (unsigned int)
- Default
  - 0 (STDIN\_FILENO): standard input
  - 1 (STDOUT\_FILENO): standard output
  - 2 (STDERR\_FILENO): standard error

# I/O System Calls

- Open files:

```
# include <fcntl.h>
```

```
int open(const char *pathname, int o_flag, ... );  
// man 2 open
```

- Return value
  - Success: file descriptor
  - Failed: -1
- o\_flag:
  - O\_RDONLY, O\_WRONLY, O\_RDWR
  - Options:
    - O\_APPEND, O\_CREAT, O\_TRUNC, ...



# I/O System Calls

- Open files
  - File descriptors: the **smallest** one available
  - Examples

```
int main (int argc, char **argv)
{
    int fd = open("foo", O_RDONLY);
    printf("%d", fd);
}
```

```
int main (int argc, char **argv)
{
    close(0);
    int fd = open("foo", O_RDONLY);
    printf("%d", fd);
}
```

# I/O System Calls

- Open files
  - STDIN\_FILENO, STDOUT\_FILENO, STDERR\_FILENO
  - opened by the OS when creating a process

# I/O System Calls

- Close files

```
# include <unistd.h>

int close(int fildes);
```

- Return

- Success: 0
- Failed: -1

# I/O System Calls

- File Position

```
# include <unistd.h>
```

```
off_t lseek(int filesdes, off_t offset, int whence);
```

- “Current file offset”:
  - An offset (in byte) to the \$whence of the file
- whence:
  - SEEK\_SET, SEEK\_CUR, SEEK\_END

# I/O System Calls

- Read files

```
# include <unistd.h>
```

```
int read(int fildes, void *buf, size_t nbytes);
```

- Start reading at “file offset”

- Return:

- Success: number of bytes read (0, if EOF)
- Failed: -1

- Return < size

- EOF
- Read from terminal (stdin), one line
- ...

# I/O System Calls

- Write files

```
# include <unistd.h>
```

```
int write(int filedes, const void *buf, size_t nbytes);
```


- Return:

- Success: number of bytes write
- Failed: -1

# An Example: I/O and Buffers

- I/O without buffer
  - No (user space) buffer
    - read(), write(): system calls
    - Do have buffer in **kernel space (by file system)**
  - Let's do some coding
- Buffering do matter!
  - printf, scanf in standard I/O library are buffered

# Revisit File Descriptors



1. Each process has its own array of "struct file\*"
2. Each file associates with only one "struct inode"
3. The "inode number" is a low-level id of a file



```

struct files_struct {
    int count;
    fd_set close_on_exec;
    fd_set open_fds;
    struct file * fd[NR_OPEN];
};

struct file {
    mode_t f_mode;
    loff_t f_pos;
    unsigned short f_flags;
    unsigned short f_count;
    unsigned long f_reada, f_ramax, f_raend, f_ralen, f_rawin;
    struct file *f_next, *f_prev;
    int f_owner;
    struct inode * f_inode;
    struct file_operations * f_op;
    unsigned long f_version;
    void *private_data;
};

struct ext2_inode {
    __u16 i_mode;        /* File type and access rights */
    __u16 i_uid;         /* Low 16 bits of Owner Uid */
    __u32 i_size;        /* Size in bytes */
    __u32 i_atime;       /* Access time */
    __u32 i_ctime;       /* Creation time */
    __u32 i_mtime;       /* Modification time */
    __u32 i_dtime;       /* Deletion Time */
    __u16 i_gid;         /* Low 16 bits of Group Id */
    __u16 i_links_count; /* Links count */
    __u32 i_blocks;      /* Blocks count */
    __u32 i_flags;       /* File flags */
    ...
    __u32 i_block[EXT2_N_BLOCKS]; /* Pointers to blocks */
    ...
};

```

# Quiz

- What happen when we open a file with a text editor?
- What happen when we open a file with two different text editors?

# A, B open the same file

Process A

fd	file pointer
0	
1	
2	
3	
4	
.	

file

file status
current offset
inode

inode

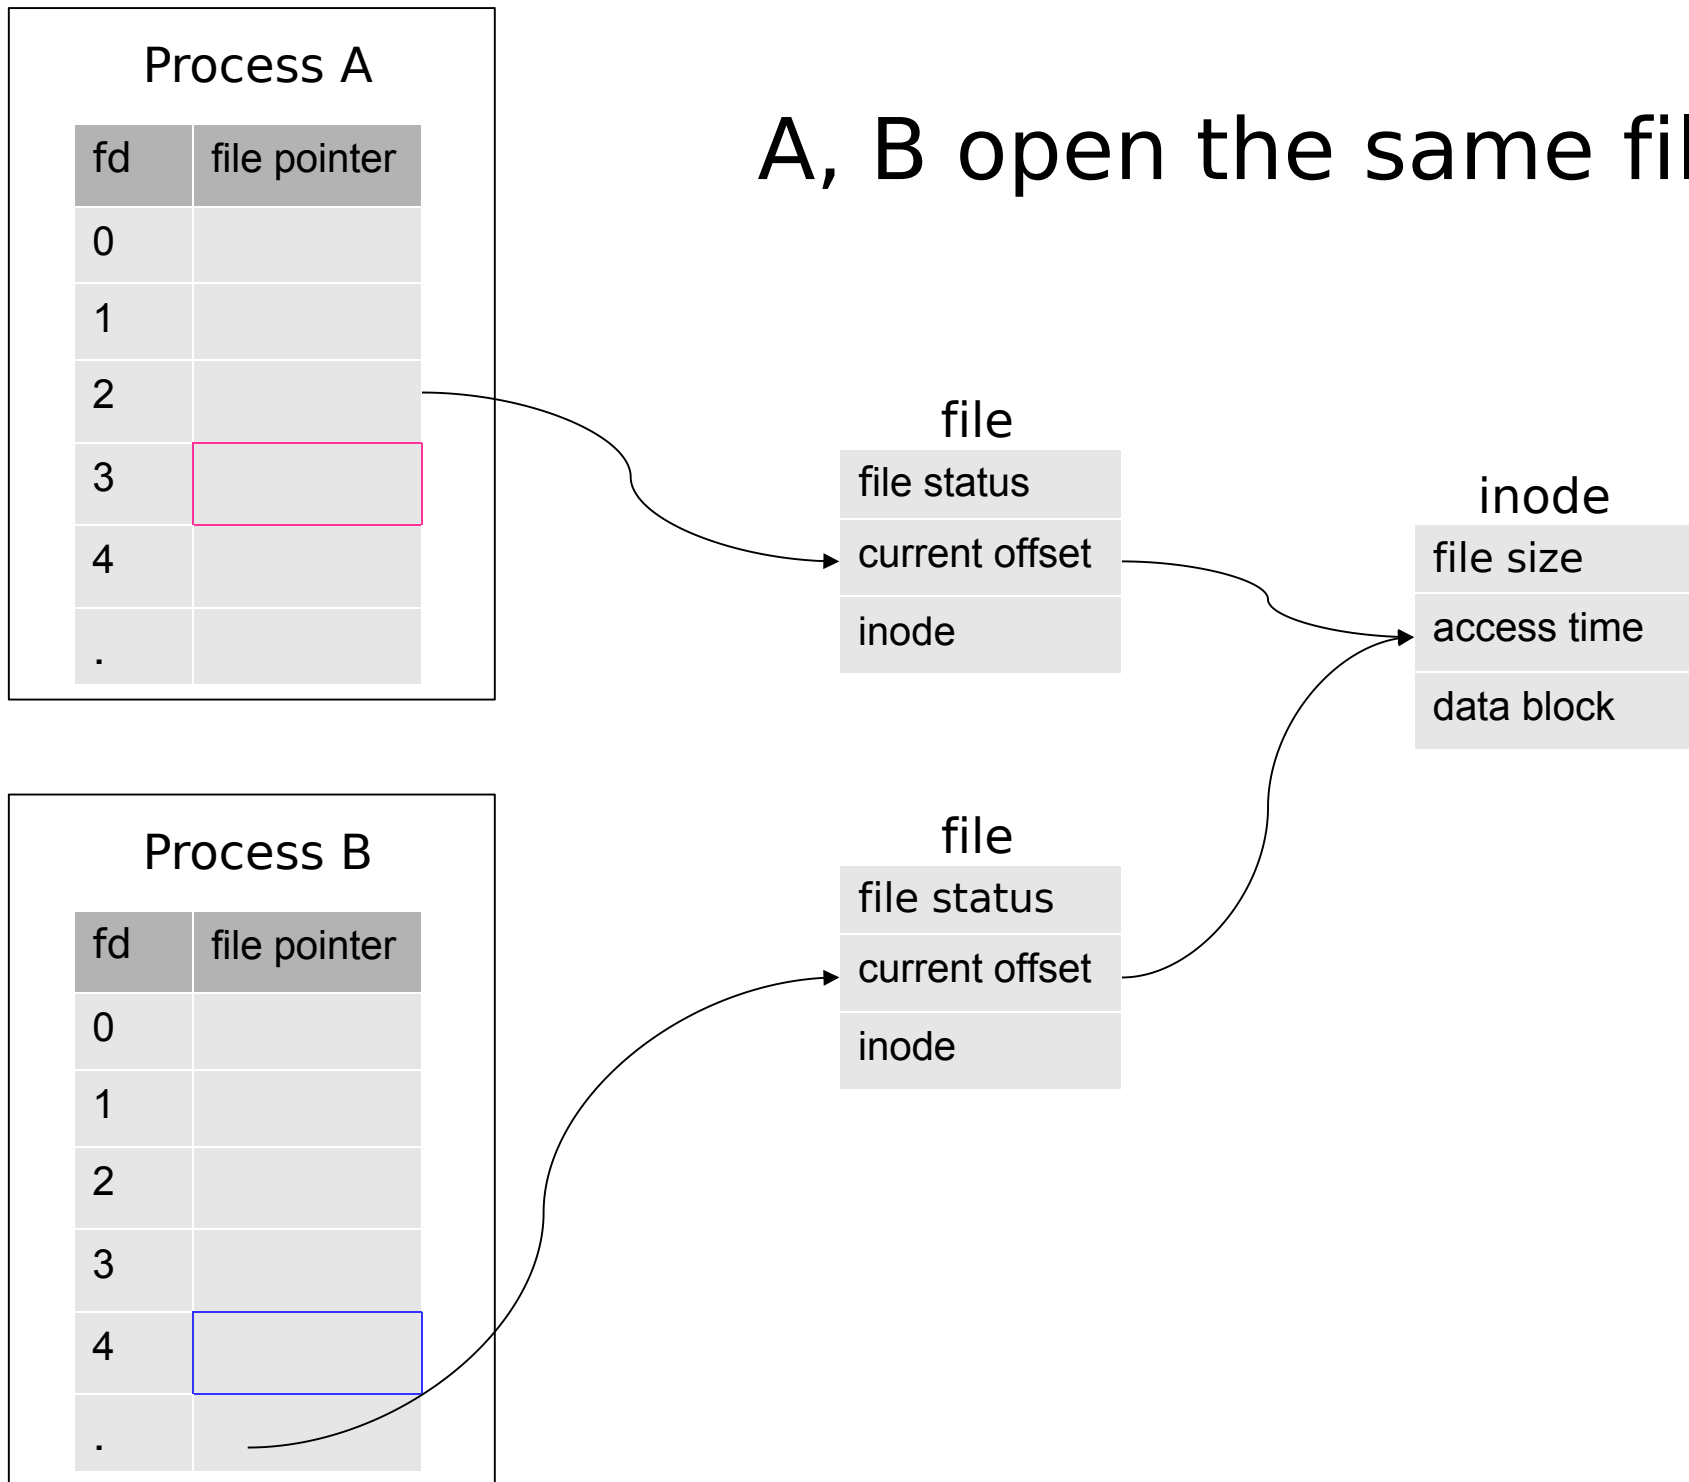
file size
access time
data block

Process B

fd	file pointer
0	
1	
2	
3	
4	
.	

file

file status
current offset
inode



# File Sharing

- Simple? ... emmm ...
- Example: how to implement
  - `open("file", O_WRONLY | O_APPEND)`
- Two process (A and B) run the same code, what will happen?

```
if (lseek(fd, 0, SEEK_END) < 0)
    perror("lseek");
```

```
if (write(fd, buf, 100) < 100)
    perror("write");
```

**Atomic operations**

# File Sharing


- Duplicate a file descriptor

```
# include <unistd.h>
```

```
int dup2(int fd, int fd2);
```

- set “fd2” point to the same file of “fd”
- Return
  - Success: fd
  - Failed: -1

```
// if fd 0 is open, close it first  
dup2(3, 0);
```



1. a file with multiple file descriptors
2. I/O redirection

# I/O System Calls

- Other system calls
  - `sync()` / `fsync()`:
    - “delay write”
    - Flush kernel buffer
  - `fcntl()`: change file (opened) attributes
  - `ioctl()`: other methods

# I/O System Calls

- Summary
  - File descriptor
  - open, close, read, write, lseek, dup
  - File sharing



# Operating System Labs

- Manipulate I/O
  - System call
    - File descriptor
    - No buffering
  - Standard library
    - FILE object
    - Buffering

# Standard I/O Library

- `#include <stdio.h>`
  - FILE object (structure)
  - Buffering
  - Formatted I/O

# System Calls vs Library Functions

- Recall:

```
#include <stdio.h>
void foo()
{
    printf("bar\n");
}
```

← User application

```
printf()
fprintf()
malloc()
atoi()
```

← Library Functions  
(Glibc)

```
write(), reads(),
mmap()
```

← System Calls

Kernel

# Standard I/O Library

```
# include <fcntl.h>
```

```
int main (int argc, char **argv)
{
    int fd = open("foo", O_RDONLY);
}
```

```
# include <stdio.h>
```

```
int main (int argc, char **argv)
{
    FILE* fp = fopen("foo", "r");
}
```

- Stream and FILE object
  - A wrapper of file descriptor
  - More information:
    - **buffer**
    - error info
    - single-byte or multi-byte

# FILE Object

- Opaque pointer
  - The implementation is hidden
  - Access the struct member through functions
- Operations on FILE objects
  - Get file descriptor: `fileno(FILE* f)`
  - Set buffer: `setbuf(FILE* f, char* buf)`

# Standard I/O Library

- Buffering
  - stdio provide a “standard I/O buffer” (**user space**)
- Three types of buffering
  - Full buffered
    - Performs I/O when the buffer is full
  - Line buffered
    - Performs I/O when encounter a newline
  - Unbuffered
    - Performs I/O immediately, no buffer

# Standard I/O Library

- Three types of buffering
  - Standard error is unbuffered
  - A stream is line buffered if it refers to terminal device, otherwise full buffered
- Write “standard I/O buffer” to disc:

```
# include <stdio.h>
```

```
int fflush(FILE *fp);
```

# Standard I/O Library

- Open/Close streams

```
# include <stdio.h>
```

```
FILE *fopen(const char* path, const char * type);
```

```
FILE *fdopen(int fd, const char * type);
```

```
int fclose(FILE* fp);
```

- Type: “r”, “w”, “a”, “r+” . . .

- Return

- Failed: NULL



# Standard I/O Library

- Character-at-a-time I/O

```
# include <stdio.h>
```

```
int getc(FILE *fp);  
int fgetc(FILE *fp);
```

```
int putc(FILE *fp);  
int fputc(FILE *fp);
```

# Standard I/O Library

- Line-at-a-time I/O

```
# include <stdio.h>
```

```
char* fgets(char *buf, int n, FILE *fp);
```

```
char* gets(char *buf);    // should never be used
```

```
int fputs(char *str, FILE *fp);
```

```
int puts(char *str);
```

# Standard I/O Library

- Direct I/O

```
# include <stdio.h>
```

```
size_t fread(void *ptr, size_t size, size_t, nobj, FILE *fp);  
size_t fwrite(void *ptr, size_t size, size_t, nobj, FILE *fp);
```

# Standard I/O Library

- Standard I/O efficiency
  - Recall: buffering in system calls
  - Let's do some coding

# Standard I/O Library

- Formatted I/O
  - printf, fprintf, scanf

# Standard I/O Library

- Summary
- `#include <stdio.h>`
  - FILE object (structure)
  - Buffering
  - Formatted I/O

# Introduction of I/O Operations

- Summary
  - System call
    - File descriptor
    - No buffering
  - Standard library
    - FILE object
    - Buffering

# Project 1

- Sorting



# Announcement

- Project 0 due
  - 21:00, Sep. 19