

# Operating System Labs

Yuanbin Wu  
cs@ecnu

# Announcement

- Project 1 due
  - 21:00 Oct. 10

# Operating System Labs

- The abstraction of process
- Process API
  - `fork()`, `exec()`, `wait()`
  - Project 1a
- CPU virtualization
  - Low level & high level mechanisms
  - Project 1b

# The Abstraction of Process

- Process
  - Running programs
- What does a running program require?
  - CPU
    - Program Counter (PC)
    - Stack Pointer / Frame Pointer
  - Memory
    - Address space
  - Disk
    - Set of file descriptors

# The Abstraction of Process

- proc file system
  - The Linux kernel has two primary functions:
    - control access to physical devices on the computer
    - schedule when and how processes interact with these devices.
  - /proc/ directory contains a hierarchy of special files
  - the current state of the kernel — allowing applications and users to peer into the kernel's view of the system.

# The Abstraction of Process

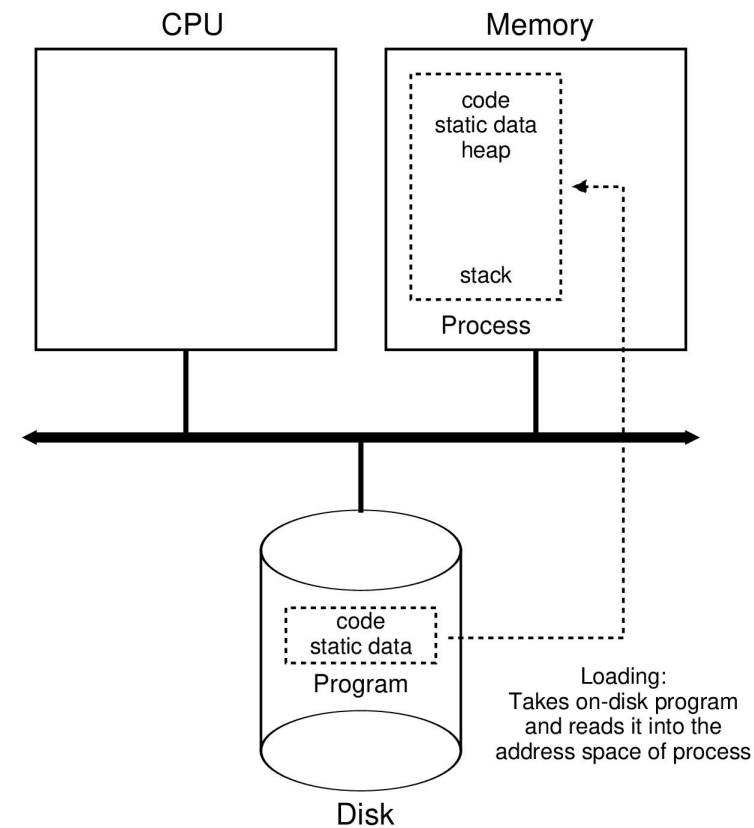
- proc file system
  - processes as files (“Everything is file”)
  - Example
    - `cat /proc/<PID>/status`
    - `cat /proc/<PID>/maps`
    - `cat /proc/<PID>/fd`
    - `cat /proc/<PID>/io`
  - `/proc/interrupts`, `/proc/meminfo`, `/proc/mounts`, `/proc/partitions`
  - Provide a method of communication between kernel space and user space
    - `ps` command

# The Abstraction of Process

- Process operations
  - Create
  - Destroy
  - Wait
  - Miscellaneous Control
  - Get status

# The Abstraction of Process


- Example: process creation
  - Load code and static data
    - local variables, function calls
  - Establish stack
    - local variables, function calls
  - Init heap
    - malloc, free
  - Allocate file descriptors
    - STDIN\_FILENO
    - STDOUT\_FILENO
    - STDERR\_FILENO





# The Abstraction of Process

- Process States



# The Abstraction of Process

- Process States

Time	Process <sub>0</sub>	Process <sub>1</sub>	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	
4	Running	Ready	Process <sub>0</sub> now done
5	–	Running	
6	–	Running	
7	–	Running	
8	–	Running	Process <sub>1</sub> now done

Time	Process <sub>0</sub>	Process <sub>1</sub>	Notes
1	Running	Ready	
2	Running	Ready	
3	Running	Ready	Process <sub>0</sub> initiates I/O
4	Blocked	Running	Process <sub>0</sub> is blocked,
5	Blocked	Running	so Process <sub>1</sub> runs
6	Blocked	Running	
7	Ready	Running	I/O done
8	Ready	Running	Process <sub>1</sub> now done
9	Running	–	
10	Running	–	Process <sub>0</sub> now done

# The Abstraction of Process

- Data structures



# Process API

- Process API
  - `fork()`, `exec()`, `wait()`, `exit()`
  - Create, execute, wait and terminate a process
  - May be the strangest API you've ever met

# Process API

- `fork()`
  - Create a new process
  - **Exactly copy** the calling process
- The return code of `fork()` is different
  - In parent: `fork()` return the pid of the child
  - In child: `fork()` return 0
- Who will run first is not determined

# Process API

- `wait()`
  - Wait for child to finish his job
  - The parent will not proceed until `wait()` return.
- `waitpid()`

# Process API

- `exec()`
  - Execute a different program in child process
- A group of system calls:
  - `execl`, `execv`, `execle`, `execve`, `execlp`, `execvp`, `fexecv`



# Process API

- Some Coding
  - fork
  - fork, wait
  - fork, wait, execvp

# Process API

- What's happening behind fork()?
  - The child get a “copy” of parent's data space, stack, heap
    - the system call: clone()
  - “Copy-on-write”
    - Not really copy the data, but share the data with “read only” flag
    - If parent or child writes on a shared address, the kernel make a copy of that piece of memory only (usually a page)

# Process API

- What's happening behind fork()?
  - File sharing
    - fd
    - File offsets

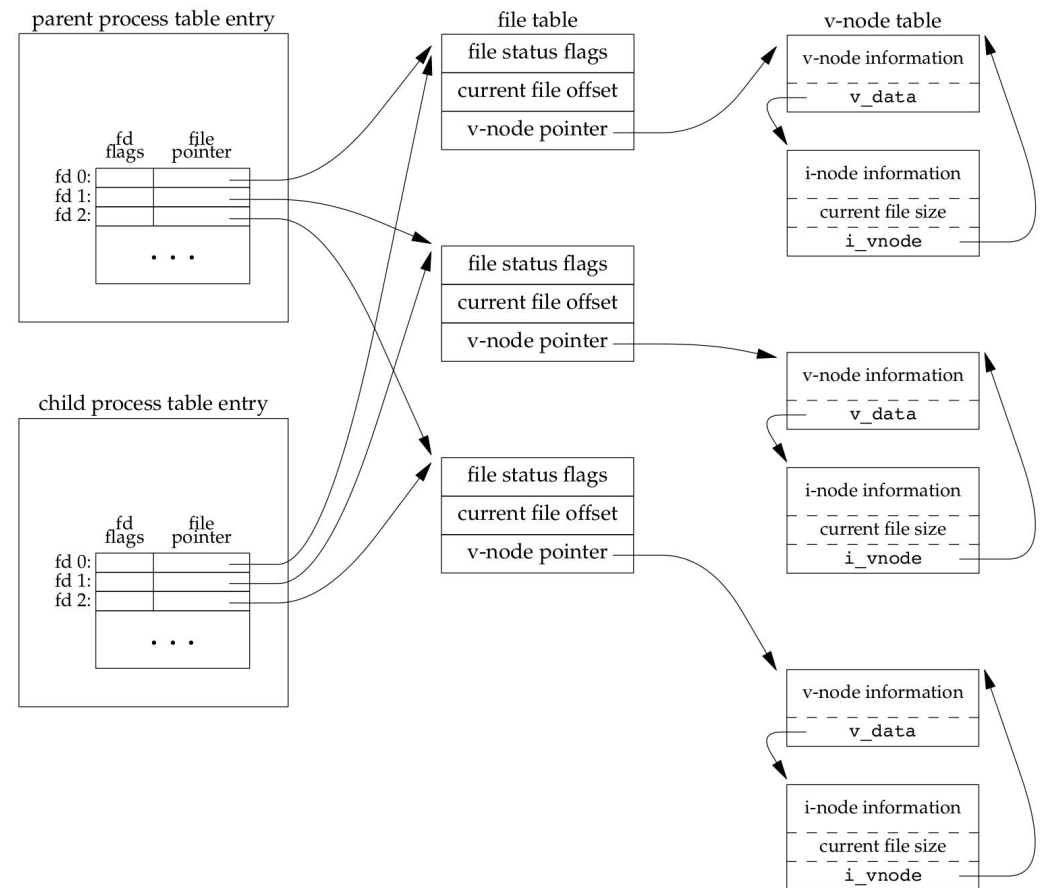


Figure 8.2 Sharing of open files between parent and child after `fork`

# Process API

- What's happening behind `fork()`?
  - Other shared data:
    - User ID, group ID...
    - Current working directory
    - Environment
    - Memory mapping
    - Resources limits
    - ...

# Process API

- What's happening behind `exit()`?
  - Close all fds, release all memory, ...
  - Inform the **exit status** to the parent process, which can be captured by `wait()`

# Process API

- What's happening behind wait()?
  - The parent terminates first?
    - The init process (PID=0)
  - The child terminates first?
    - The kernel keeps a small amount of information for every terminating process
    - Available when the parent calls wait()
      - PID, termination status the amount of CPU time
    - zombies




# Process API

- What's happen behind wait()/waitpid()
  - wait(): block the caller until a child process terminates
  - waitpid(): wait which child, and some other options

# Process API

- What's happening behind `exec()`?
  - Replace the current process with a new program from disk
    - Text, data, heap, stack
  - Start from the `main()` of that program





# Process API

- Process API summary
  - `fork()`: create a new process
  - `wait()`: wait for a child
  - `exit()`: destroy a process
  - `exec()`: execute a program in child

# Project1a

- Implement your own shell
  - Use fork, wait, execvp
  - Also open, close, dup2

# Project1a Details

- Basic shell
  - Run your shell by: `./mysh`
  - It will print a prompt:

```
mysh>
```

- You can type some commands

```
mysh> ls
```

- Hit ENTER, the command will be executed

# Project1a Details

- Build-in Commands
  - When “mysh” execute a command, it will check weather it is a **build-in** or not.
  - For build-in commands, you should involve your implementation.
  - They are:
    - exit
    - wait
    - cd
    - pwd

# Project1a Details

- Redirection

- Your shell should support redirection:

```
mysh> ls -l > output
```

- The file “output” contain the result of “ls -l”

# Project1a Details

- Background Jobs

- Your shell should be able to run jobs in the background

mysh> ls &

- Your shell will continue to work rather than wait.

# Project1a Details

- Batch mode

- Your shell should be able to run in batch mode

`./mysh batch_file`

- Your shell will run the commands in batch\_file
- E.g, “batch\_file” contains
  - `ls -l`
  - `echo hello`

# Project1a Details

- Bonus: Pipe

- The pipe connect the input/output of different commands

```
mysh> grep "hello" FILE | wc -l
```

- How many lines have "hello"



# CPU Virtualization

- What
  - Provide the illusion of many CPUs
- Why
  - Multi-task
- How
  - Time sharing

# CPU Virtualization

- Mechanisms
  - Low level mechanisms
    - Context switch
  - High level intelligence
    - Scheduling policy


# CPU Virtualization

- Low level mechanisms
  - Direct Execution
    - Just run a program on CPU directly

OS	Program
Create entry for process list	
Allocate memory for program	
Load program into memory	
Set up stack with argc/argv	
Clear registers	
Execute <b>call</b> main()	Run main()
	Execute <b>return</b> from main
Free memory of process	
Remove from process list	

# Direct Execution

- Problems of direct execution
  - Visit any memory address
  - Open any file
  - Directly play with hardwares (e.g. I/O)



**Lost control**

# Limited Direct Execution

- Limited Direct Execution
  - Kernel model and user model
  - “restricted operations”
    - By OS
  - When a thread needs “restricted operations”
    - System call

# Limited Direct Execution

- User mode
  - The behavior of the code is restricted
- Kernel mode
  - The code can do what it likes to do
    - Issue I/O, executing all types of instructions,...
- How to switch?
  - System call

# System Call

- Hardware supports on system call
  - A bit in CPU identifies kernel/user mode
  - “trap” instruction
  - “return-from-trap” instruction
  - Save the registers before do the restricted operation (kernel stack)

OS @ run  
(kernel mode)

Hardware

Program  
(user mode)

---

Run main()

...

Call system call  
**trap** into OS



OS @ run  
(kernel mode)

Hardware

Program  
(user mode)

---

Run main()

...

Call system call  
**trap** into OS

save regs to kernel stack  
move to kernel mode  
jump to trap handler

OS @ run  
(kernel mode)

Hardware

Program  
(user mode)

---

Run main()

...  
Call system call  
**trap** into OS

save regs to kernel stack  
move to kernel mode  
jump to trap handler

Handle trap  
Do work of syscall  
**return-from-trap**

OS @ run  
(kernel mode)

Hardware

Program  
(user mode)

---

Run main()  
...  
Call system call  
**trap** into OS

save regs to kernel stack  
move to kernel mode  
jump to trap handler

Handle trap  
Do work of syscall  
**return-from-trap**

restore regs from kernel stack  
move to user mode  
jump to PC after trap

**OS @ run  
(kernel mode)**

**Hardware**

**Program  
(user mode)**

---

Create entry for process list  
Allocate memory for program  
Load program into memory  
Setup user stack with argv  
Fill kernel stack with reg/PC  
**return-from-trap**

restore regs from kernel stack  
move to user mode  
jump to main

Run main()  
...  
Call system call  
**trap** into OS

save regs to kernel stack  
move to kernel mode  
jump to trap handler

Handle trap  
Do work of syscall  
**return-from-trap**

restore regs from kernel stack  
move to user mode  
jump to PC after trap

**OS @ run  
(kernel mode)**

**Hardware**

**Program  
(user mode)**

Create entry for process list  
Allocate memory for program  
Load program into memory  
Setup user stack with argv  
Fill kernel stack with reg/PC  
**return-from-trap**

restore regs from kernel stack  
move to user mode  
jump to main

Run main()  
...  
Call system call  
**trap** into OS

save regs to kernel stack  
move to kernel mode  
jump to trap handler

Handle trap  
Do work of syscall  
**return-from-trap**

restore regs from kernel stack  
move to user mode  
jump to PC after trap

...  
return from main  
**trap** (via `exit()`)

Free memory of process  
Remove from process list

**OS @ boot  
(kernel mode)**

**Hardware**

initialize trap table

remember address of...  
syscall handler

**OS @ run  
(kernel mode)**

**Hardware**

**Program  
(user mode)**

Create entry for process list  
Allocate memory for program  
Load program into memory  
Setup user stack with argv  
Fill kernel stack with reg/PC  
**return-from-trap**

restore regs from kernel stack  
move to user mode  
jump to main

Run main()  
...  
Call system call  
**trap** into OS

save regs to kernel stack  
move to kernel mode  
jump to trap handler

Handle trap  
Do work of syscall  
**return-from-trap**

restore regs from kernel stack  
move to user mode  
jump to PC after trap

...  
return from main  
**trap** (via `exit()`)

Free memory of process  
Remove from process list

# Limited Direct Execution

- Switching between processes
  - Cooperative approach
    - OS trusts the process to yield CPU properly
  - Non-cooperative approach
    - OS revokes the control of CPU periodically
    - Time interrupt
    - Scheduler

OS @ boot (kernel mode)	Hardware	
initialize trap table	remember addresses of... syscall handler timer handler	
start interrupt timer	start timer interrupt CPU in X ms	
OS @ run (kernel mode)	Hardware	Program (user mode)
		Process A
		...
	<b>timer interrupt</b> save regs(A) to k-stack(A) move to kernel mode jump to trap handler	
Handle the trap		
Call <code>switch()</code> routine		
save regs(A) to proc-struct(A)		
restore regs(B) from proc-struct(B)		
switch to k-stack(B)		
<b>return-from-trap (into B)</b>		
	restore regs(B) from k-stack(B) move to user mode jump to B's PC	
		Process B
		...



# Limited Direct Execution

- Low-level mechanisms: summary
  - Direct execution
  - Limited direct execution
  - Switch between processes

# •Scheduling Policy

- High level intelligence
  - Scheduling policy
    - First In, First Out
    - Shortest job first
    - Shortest time to complete first
    - Round Robin

# •CPU virtualization

- Summary of CPU virtualization
  - Low level mechanisms
    - A little hardware support goes a long way
  - High level mechanisms

# Project1b Details

- Adding a system call for xv6
  - Understanding the low-level mechanism
  - Kernel mode, user mode
  - Trap
  - Interrupt handler

# Project1b Details

- The system call
  - `int getreadcount()`
  - Return how many times the `read()` system call has been called

# Project1b Details

- Get familiar with xv6
  - QEMU emulator
    - Installed with make
  - Compile and run xv6
    - Compile: make
    - Run: make qemu-nox
    - Debug: make qemu-nox-gdb