

Operating System Labs

Yuanbin Wu
CS@ECNU

Operating System Labs

- Project 2 Due
 - 21:00, Nov. 6
- Project 2
 - group of 3
 - You now have 3 “late days”, but start early!
 - Oral test at week 12 (Nov. 22)
- We encourage using **git** for your group assignments (geeks' choice).

Operating System Labs

- Update of timeline
 - Week 5 (Oct. 3) --> Oct. 8 (Today)
 - Week 6 (Oct. 10) lecture
 - Week 7 (Oct. 17) no lecture

Operating System Labs

- C Memory API
- Free Memory Management

C Memory API

- Type of memory
 - Stack
 - Heap

C Memory API

- Stack
 - Allocated / Deallocate automatically
 - By the compiler
 - Automatic memory

C Memory API

- Stack

- Example (local variable)

```
void func()
{
    int x = 0;
    ...
}
```

- You only declare the variable
- Compiler will allocate it when someone calls the function
- Also will deallocate it when func returns

C Memory API

- Heap
 - Allocated / Deallocate explicitly
 - By you, the programmer

C Memory API

- Heap

- Example (malloc)

```
void func()
{
    int *ptr = (int*)malloc(sizeof(int));
    ...
}
```

- Both stack and heap allocation
 - When func returns,
 - Stack memory will be deallocated
 - Heap memory is still there

C Memory API

- Stack and Heap

- Heap

- From low addr to high addr


- Stack

- From high addr to low addr

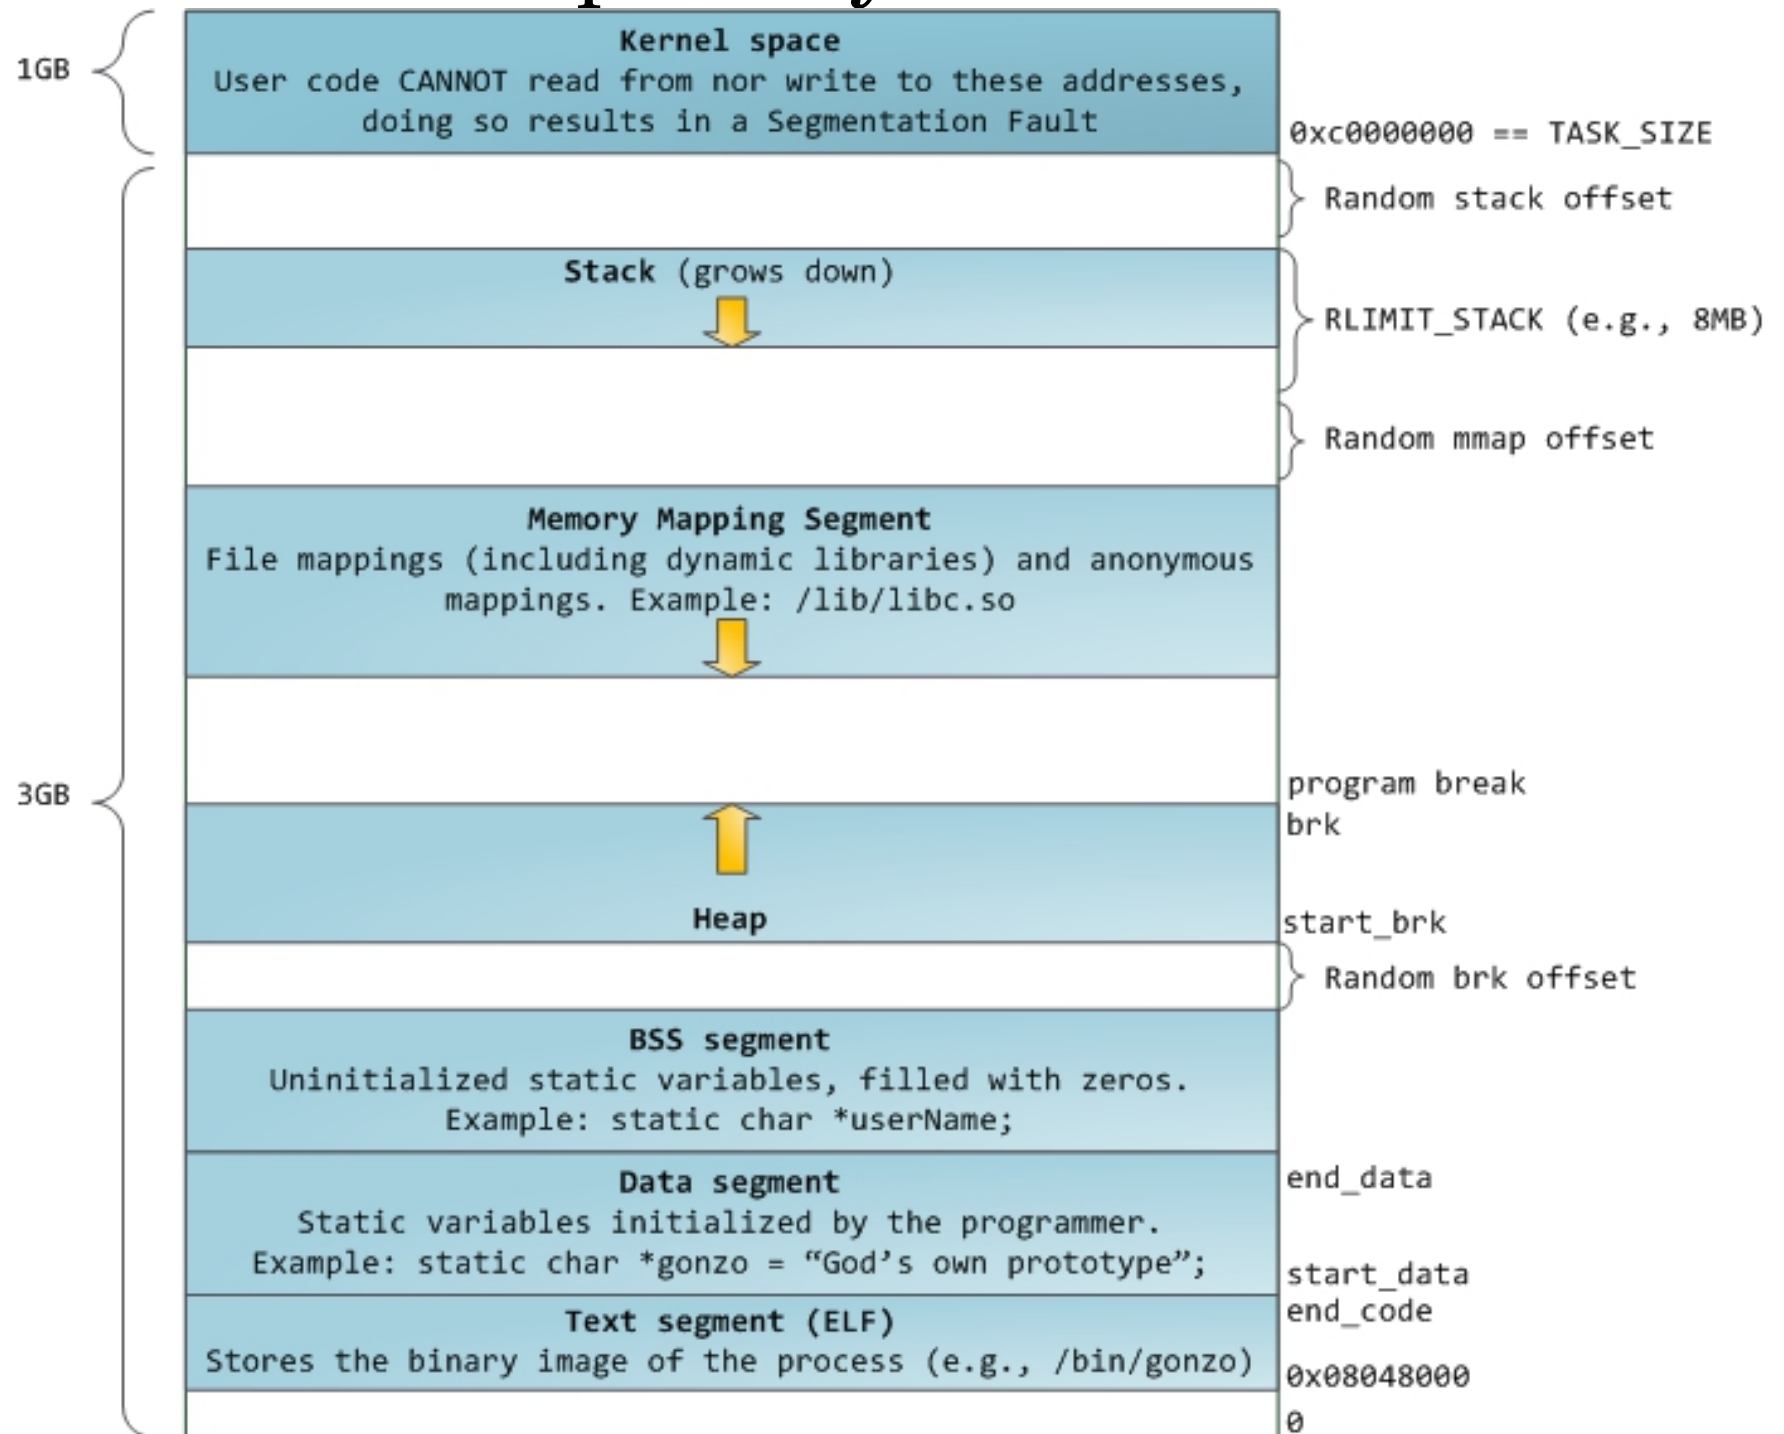
- Let's see

- mmap

- front-end for smap file proc file system



• A real address space layout



C Memory API

- Malloc


```
#include <stdlib.h>  
void *malloc(size_t size);
```

- If failed, return NULL

- Free


```
#include <stdlib.h>  
void free(void* ptr);
```

Allocation





-----> pointer

```
int *pi; // local variable
```



```
pi = (int *)malloc(sizeof(int) *  
4);
```

Free



C Memory API

- Segment fault

```
char *src = "hello";  
char *dst;  
strcpy(dst, src);
```

- run this code, it will likely lead to a

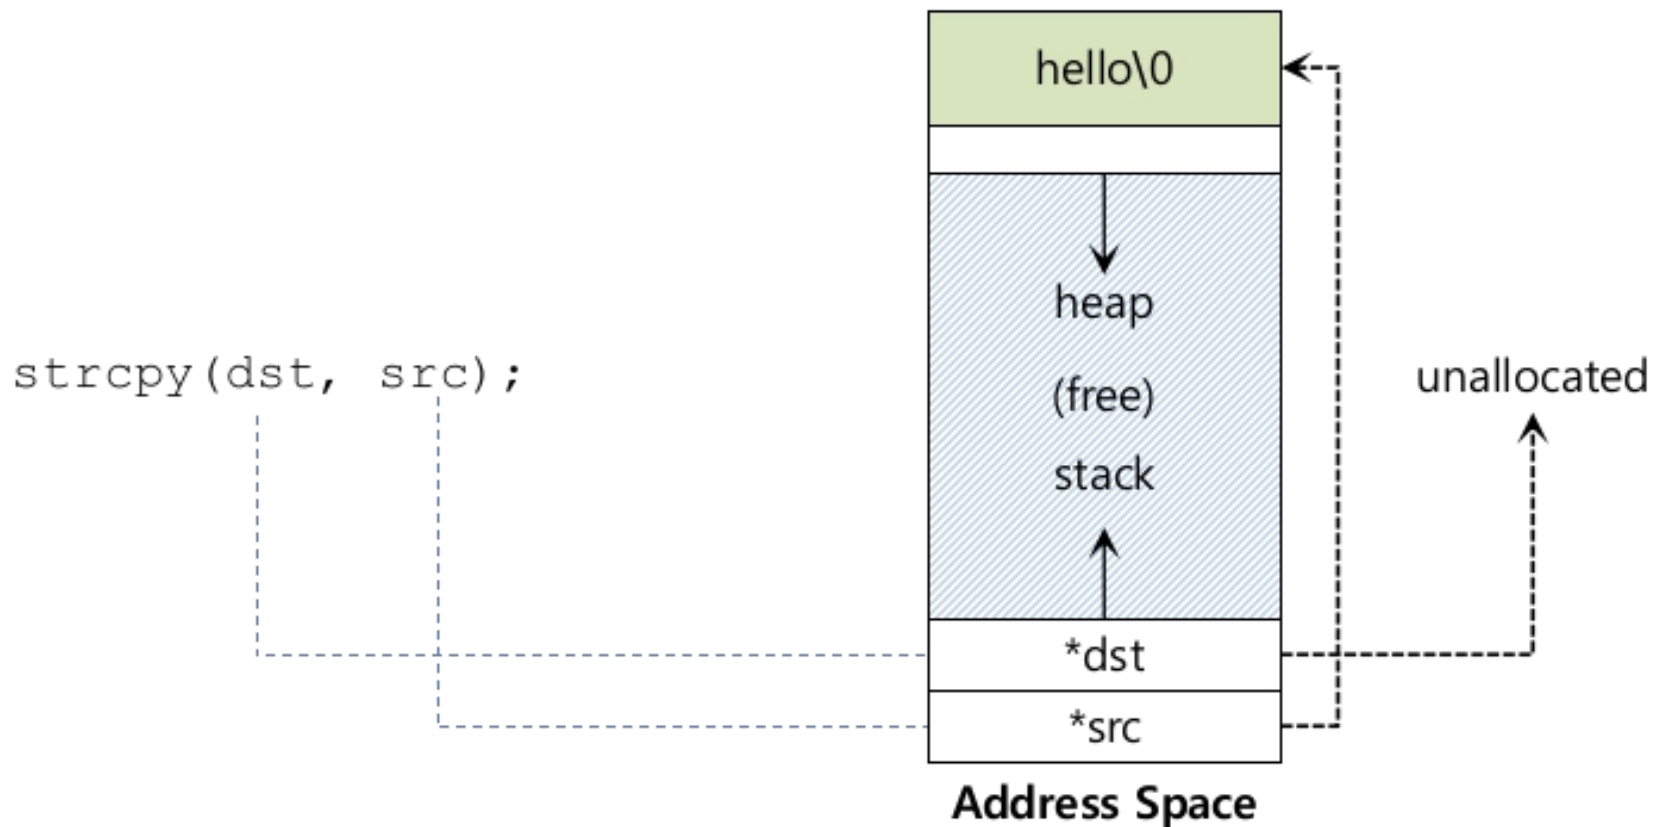
segmentation fault

- It is a fancy term for

YOU DID SOMETHING WRONG WITH MEMORY
YOU FOOLISH PROGRAMMER AND I AM ANGRY.


Segmentation Fault

```
char *src = "hello"; //character string constant  
char *dst;           //unallocated  
strcpy(dst, src);    //segfault and die
```



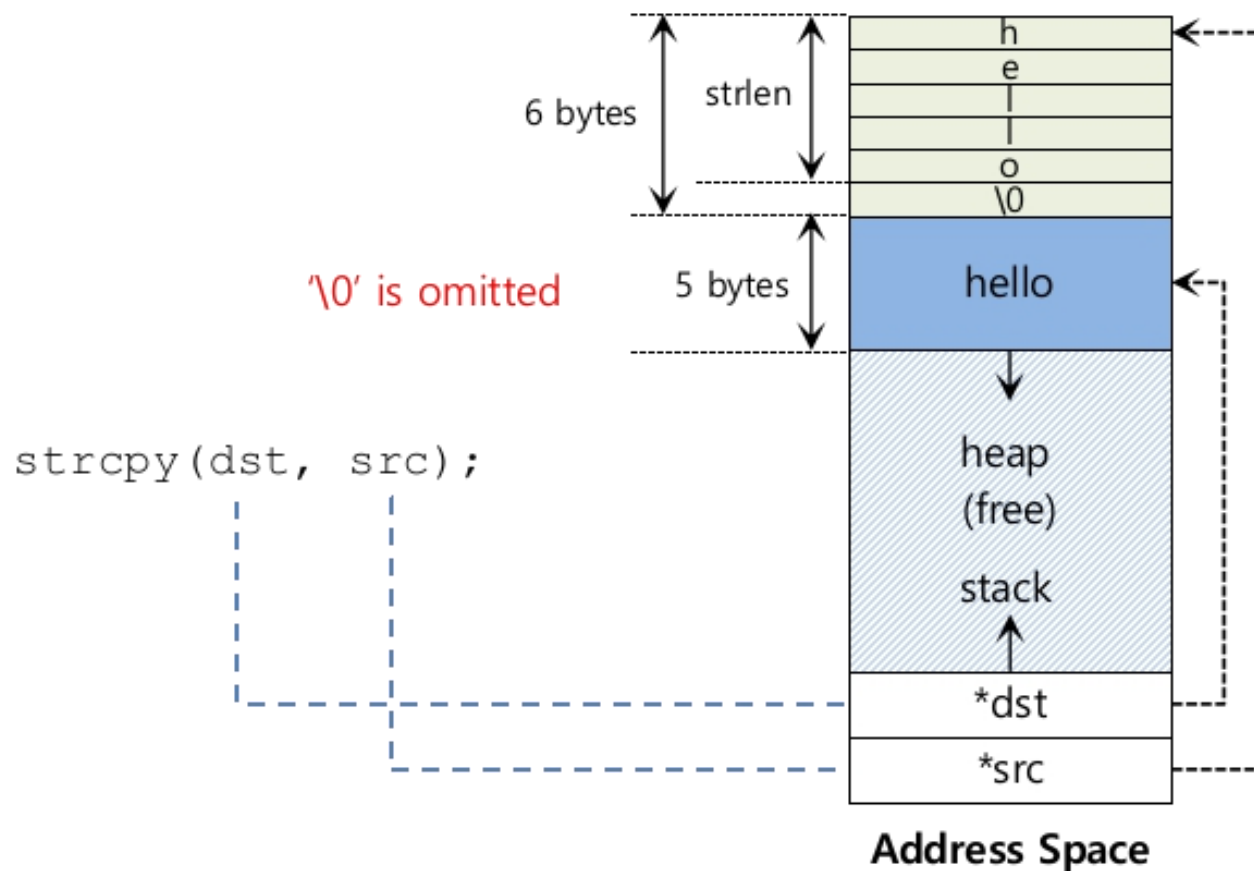
Correct Code

```
char *src = "hello";    //character string constant
char *dst (char *)malloc(strlen(src) + 1 ); // allocated
strcpy(dst, src);        //work properly
```



Works, but buggy


```
char *src = "hello"; //character string constant
char *dst (char *)malloc(strlen(src)); // too small
strcpy(dst, src);    //work properly
```




Uninitialized Read

- Wild pointer

```
int *x = (int *)malloc(sizeof(int)); // allocated
printf("*x = %d\n", *x); // uninitialized memory access
```




Address Space




Address Space

Memory Leak


unused : unused, but not freed



Address Space






Address Space




Address Space

run out of memory

Dangling Pointer



`free`
(b) →



C Memory API

- Standard library
 - malloc(), realloc(), free()
- System calls
 - brk(), sbrk()
 - mmap()
- For comparison
 - printf() and write()
 - “Buffer the system call”

System calls: brk(), sbrk()

```
#include <unistd.h>
```

```
int brk(void *addr)
```


```
void *sbrk(intptr_t increment);
```

- brk/sbrk: expand the program's break.
 - break: The location of the end of the heap in address space

System calls: mmap()

```
#include <sys/mman.h>
```

```
void *mmap(void *ptr, size_t length, int prot, int flags,  
int fd, off_t offset)
```




C Memory API

- Summary: common errors
 - Forget to allocate memory
 - Not allocating enough memory
 - Forget to initialize allocated memory
 - Forget to free memory
 - Free memory before you are done with it
 - Free memory repeatedly
 - Call free() incorrectly



AND SUDDENLY YOU MISSTEP, STUMBLE, AND JOLT AWAKE?



Free Memory Management




Dark Forest of Pointers

Free Memory Management

- Manage a bunch of “memory”
- Where ?
 - implementation of virtual memory (kernel space)
 - implementation of memory allocation libs (user space)

Free Memory Management

- Fixed-size unit
 - Paging
 - Problem: internal fragmentation
- Variable-size unit
 - User level: memory allocation library
 - Kernel level: VM implemented with segmentation
 - Problem: external fragmentation



Free Memory Management

- Free memory management
 - How to manage variable-size free memory units
 - How to implement
 - `malloc(size_t size)`
 - `free(void *ptr)`

Free Memory Management



- Assumptions
 - Focus on external fragmentation
 - No compaction
 - Manage a contiguous region of bytes (by `mmap()` system call)

Free Memory Management

- Low-level Mechanisms
 - Splitting and Coalescing
 - Tracking allocated regions
 - Implementation of a free list
- High-level Intelligence
 - Best fit
 - Worst fit
 - First fit
 - Next fit

Free Memory Management

- **Splitting and Coalescing**
 - Free list: a set of free chunks
 - Two chunks (10 bytes each)





Free Memory Management

- **Splitting and Coalescing**

- request less than 10 bytes? (e.g. malloc(1))




- **Splitting**



Free Memory Management

- **Splitting and Coalescing**

- Free a chunk?



- Malloc(20)?

- **Coalescing**



Free Memory Management

- **Tracking Allocated Regions**

- Observation on `free(void *ptr)`
 - No size parameter
- Given a pointer, the malloc library could determine the size of region
- How?
 - Some extra information
 - **header** of a memory block


Free Memory Management

- **Tracking Allocated Regions**

- header


```
typedef struct __header_t {  
    int size;  
    int magic;  
} header_t;
```

- malloc(20)



Free Memory Management

- **Tracking Allocated Regions**
 - header: example



Free Memory Management

- **Tracking Allocated Regions**

- free(ptr)

- Get the size of the region

```
void free(void *ptr) {  
    header_t *hptr = (void *)ptr - sizeof(header_t);  
}
```


- Check whether ptr is valid

```
assert(hptr->magic == 1234567)
```

Free Memory Management

- **Implementation of the Free List**

- Free list



- Implementation

- List node (allocate a node when needed)
 - Can NOT do this here! All you have is a given free space
 - How to build a free list **inside** the free space?

Free Memory Management


- **Implementation of the Free List**
 - Node in free list

```
typedef struct __node_t {  
    int size;  
    struct __node_t *next;  
} node_t;
```

Free Memory Management

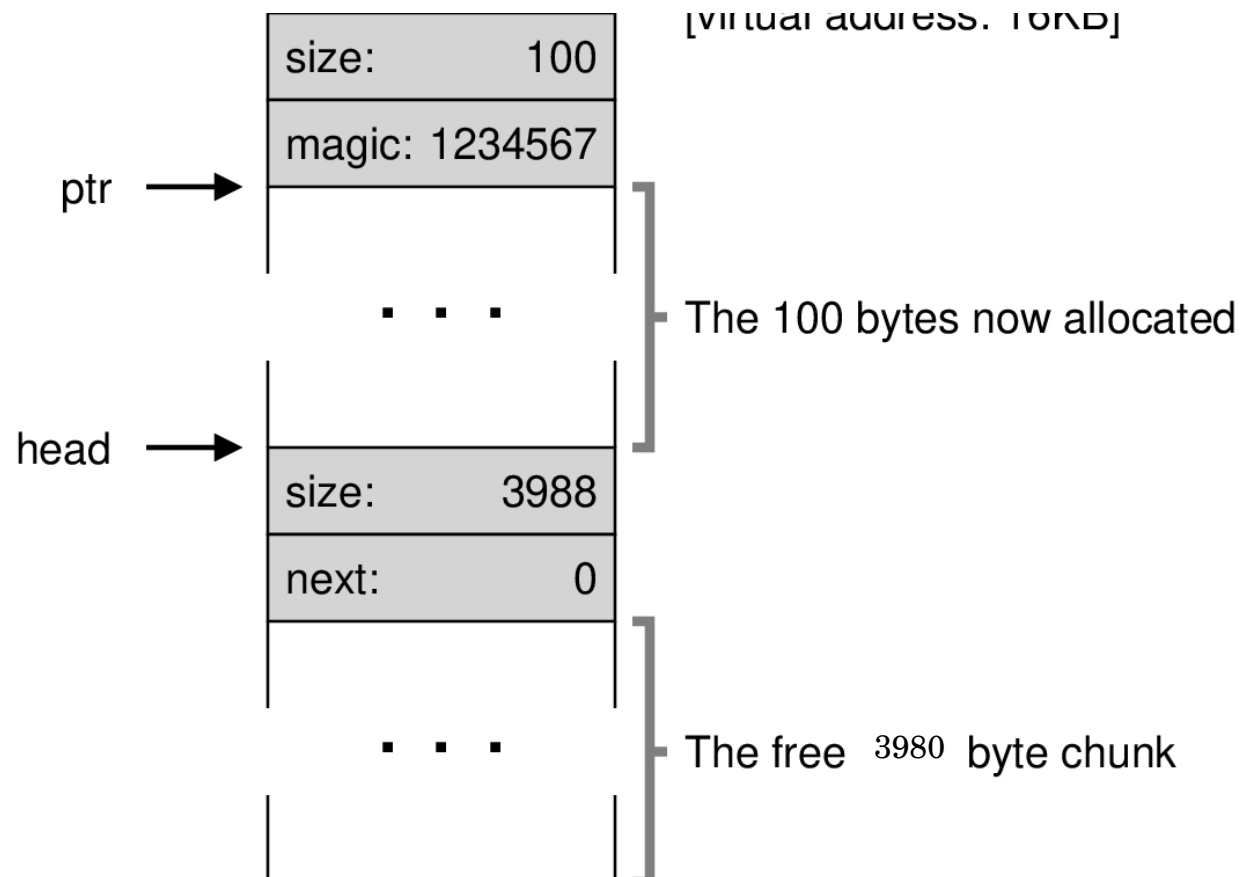
- **Implementation of the Free List**
 - Initialization (e.g. 4096)

```
// mmap() returns a pointer to a chunk of free space
node_t *head = mmap(NULL, 4096, PROT_READ |
                    PROT_WRITE, MAP_ANON | MAP_PRIVATE, -1, 0);
head->size = 4096 - sizeof(node_t);
head->next = NULL;
```

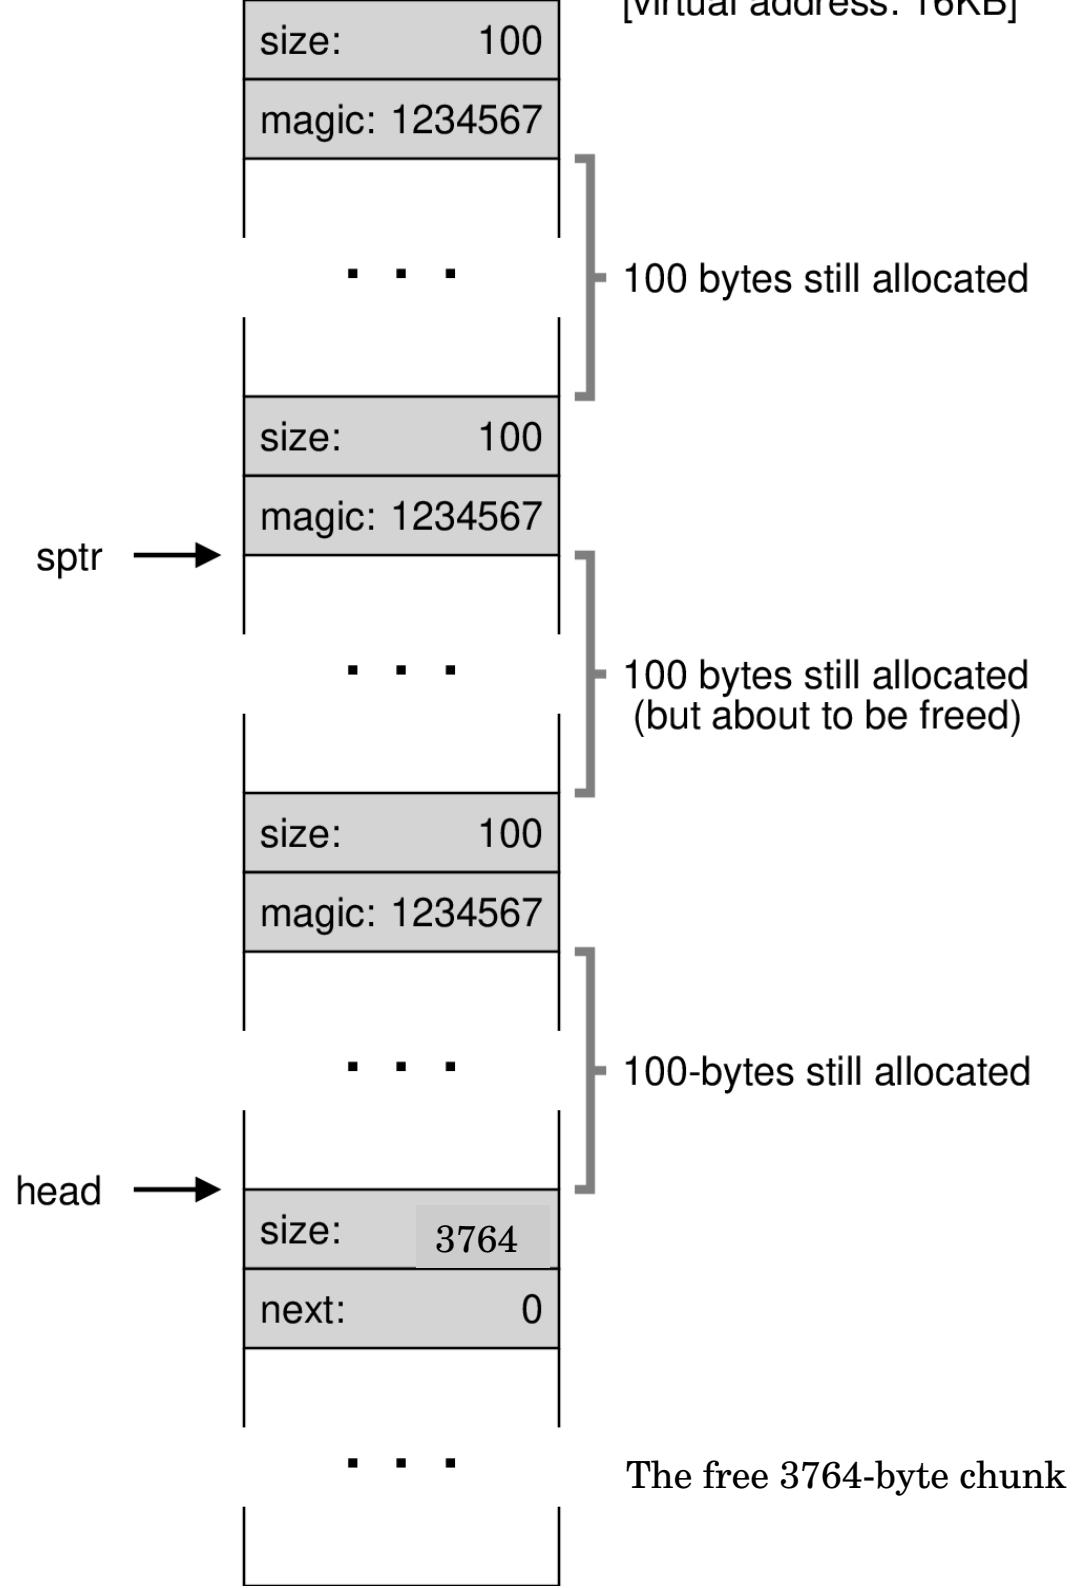


Free Memory Management

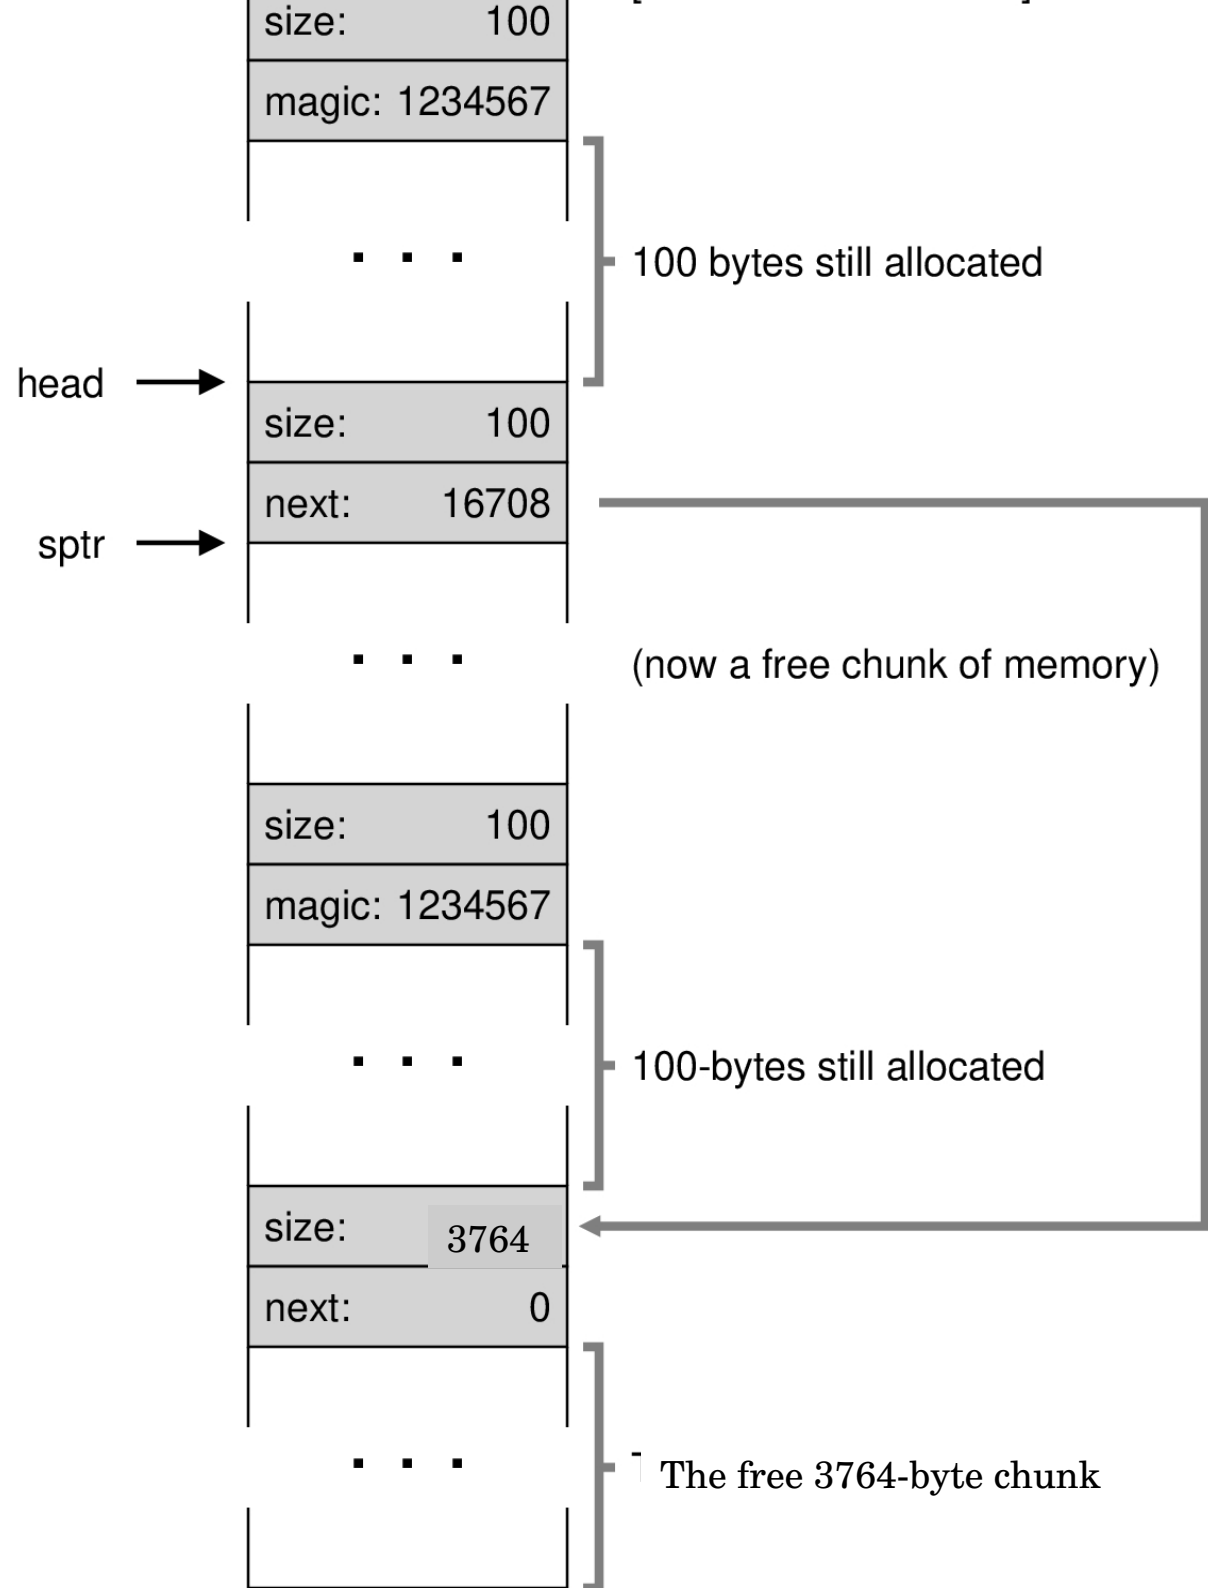
- **Implementation of the Free List**
 - `malloc(100)`



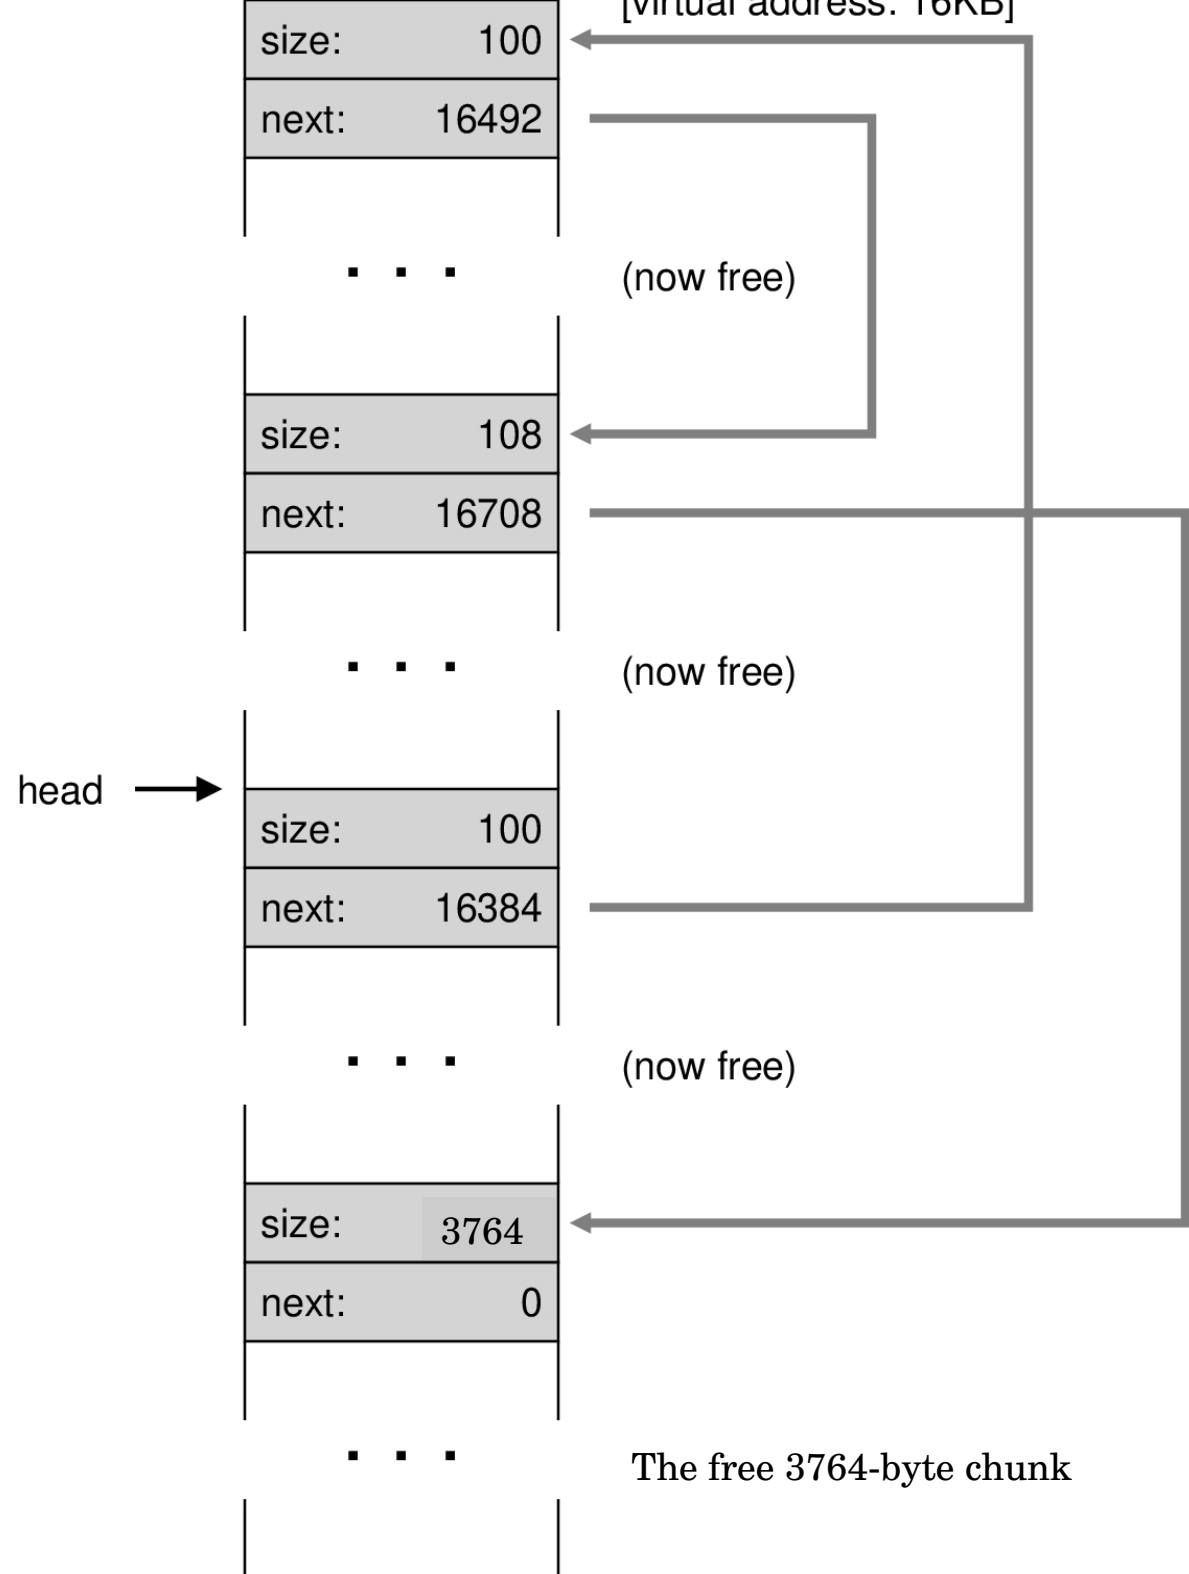
- `malloc(100)*3`



- **Free(16500)**
– $16384 + 108 + 8$



- Free()*3
- Coalesce
 - Merge adjacent
 - chunks



Free Memory Management

- Growing the Heap
 - What if the heap runs out of space?
 - Return NULL
 - Increase the size of heap
 - OS find free physical pages
 - Map them into address space of the process

Free Memory Management

- Summary of low-level Mechanisms
 - Splitting and Coalescing
 - Tracking allocated regions
 - Implementation of a free list
 - Growing the heap

Free Memory Management


- High-level intelligence
 - How to find the proper nodes in the free list?
 - Less fragmentation
 - Fast allocation
 - Some simple strategies
 - The stream of allocation and free requests can be arbitrary
 - Any strategy could be arbitrarily bad/good

Free Memory Management

- Best Fit
 - Find the smallest feasible node
- Worst Fit
 - Find the largest feasible node
- First Fit
 - Find the first feasible node

Free Memory Management

- Example



- Best fit




- Worst fit



Free Memory Management

- Other approaches
 - Segregated List
 - Slab allocator
 - Buddy Allocation
 - Binary search tree



Free Memory Management

- Dlmalloc (Doug Lea allocator)
 - Segregated list for small size allocations
 - Search the free list
 - sbrk and mmap
 - ...
- <http://g.oswego.edu/dl/html/malloc.html>
- https://cs61.seas.harvard.edu/wiki/images/e/e2/Lec11-Dynamic_memory_2.pdf